



Final project for the Deep Learning course

Da Vinchie Lisa

SIMPLE IMAGE INPAINTING

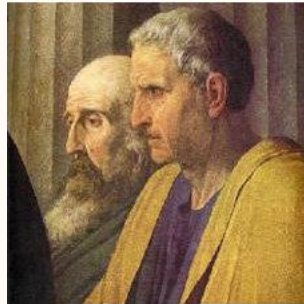
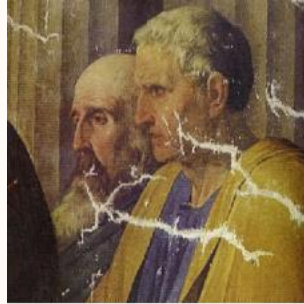


01

Introduction

What is image inpainting?

- It is a process used to restore missing parts of an image by using information from the surrounding pixels
- Examples: restoring scratched old photos, removing objects from a picture



What is image inpainting?

- It is a process used to re missing parts of an image using information from the surrounding pixels
- Examples: restoring scratched old photos, removing objects from a picture



What is image inpainting?

Challenges:

- There isn't a single correct way to fill in a missing part of an image



Main Techniques

- Partial differential equations
- Patch based
- Deep-learning based

Deep-Learning based techniques

- Autoencoders
- GANs
- Transformers
- Diffusion Models

Why the autoencoder?

Training stability

- Deterministic loss
- No adversarial training

Why the autoencoder?

Computational efficiency

- Single forward pass
- No discriminator => lower memory usage
- Faster training than GANs and diffusion models

Disadvantages

- Less realistic results than GANs
- Deterministic result
- Less accurate textures

The dataset

Cifar10

- 50000 training images
- 10000 validation images
- 3 x 32 x 32 images
- 10 classes



The dataset

Cifar10

- 50000 training images
- 10000 validation images
- 3 x 32 x 32 images
- 10 classes



The dataset

Cifar10

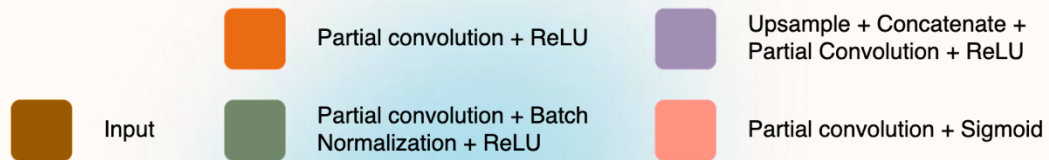
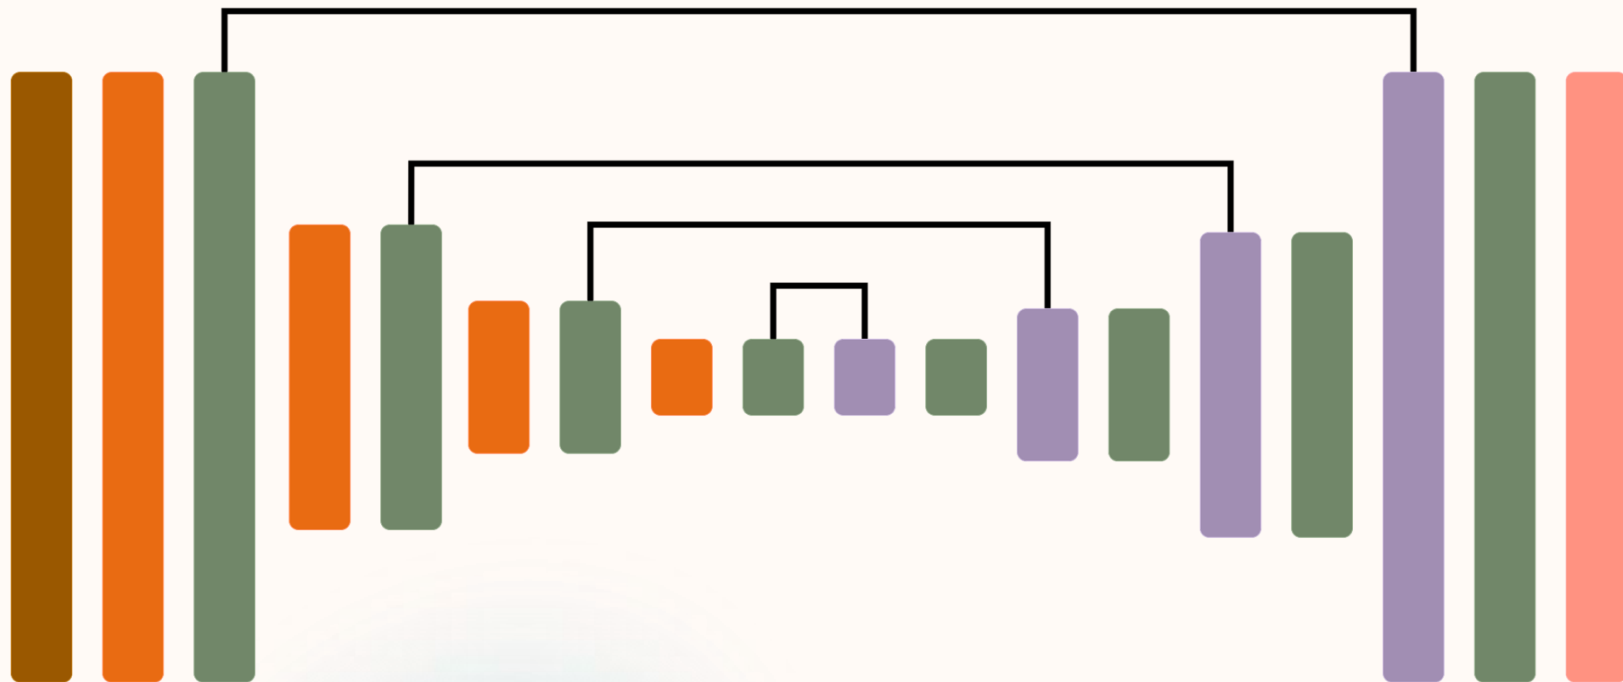
- 50000 training images
- 10000 validation images
- 3 x 32 x 32 images
- 10 classes





02

Architecture



Encoder

#	Layer	Output Size	#	Layer	Output Size
1	<u>Input</u>	3 x 32 x 32	8	<u>Partial Conv + ReLU</u>	128 x 8 x 8
2	<u>Partial Conv + ReLU</u>	32 x 32 x 32	9	<u>Partial Convolution</u>	128 x 4 x 4
3	<u>Partial Convolution</u>	32 x 16 x 16	10	<u>Batch Normalization + ReLU</u>	128 x 4 x 4
4	<u>Batch Normalization + ReLU</u>	32 x 16 x 16	11	<u>Partial Conv + ReLU</u>	256 x 4 x 4
5	<u>Partial Conv + ReLU</u>	64 x 16 x 16	12	<u>Partial Convolution</u>	256 x 2 x 2
6	<u>Partial Convolution</u>	64 x 8 x 8	13	<u>Batch Normalization + ReLU</u>	256 x 2 x 2
7	<u>Batch Normalization + ReLU</u>	128 x 8 x 8			

Decoder

#	Layer	Output Size	#	Layer	Output Size
14	<u>Upsample</u>		24	<u>Upsample</u>	
15	<u>Concatenate</u>		25	<u>Concatenate</u>	
16	<u>Partial Convolution + ReLu</u>	256 x 4 x 4	26	<u>Partial Convolution + ReLu</u>	64 x 16 x 16
17	<u>Partial Convolution</u>	128 x 4 x 4	27	<u>Partial Convolution</u>	32 x 16 x 16
18	<u>Batch Normalization + ReLU</u>	128 x 4 x 4	28	<u>Batch Normalization + ReLU</u>	32 x 16 x 16
19	<u>Upsample</u>		29	<u>Upsample</u>	
20	<u>Concatenate</u>		30	<u>Concatenate</u>	
21	<u>Partial Convolution + ReLu</u>	128 x 8 x 8	31	<u>Partial Convolution + ReLu</u>	32 x 32 x 32
22	<u>Partial Convolution</u>	64 x 8 x 8	32	<u>Partial Convolution</u>	3 x 32 x 32
23	<u>Batch Normalization + ReLU</u>	64 x 8 x 8	33	<u>Batch Normalization + ReLU</u>	3 x 32 x 32
			34	<u>Partial Convolution + Sigmoid</u>	3 x 32 x 32

Why this architecture?

- Based on previous works, adapted for smaller images
- Convolutions are very efficient in learning features and are translation invariant
- Batch Normalization helps with convergence
- Skip connections mitigate information loss during
- The sigmoid maps the values of the pixels to values between 0 and 1
- Hierarchical layers learn different features at each layer

Training an autoencoder

Training

- Take a fully known image
- Simulate a mask by replacing values under a specific area with a placeholder
- During training, reconstruct the whole image but calculate the loss only on the masked pixels

Convolutional layer

$$x' = \sum (W^T \circ X)$$

Convolutional layer

7	2	3	3	8
4	5	3	8	4
3	3	2	8	4
2	8	7	2	7
5	4	4	5	4

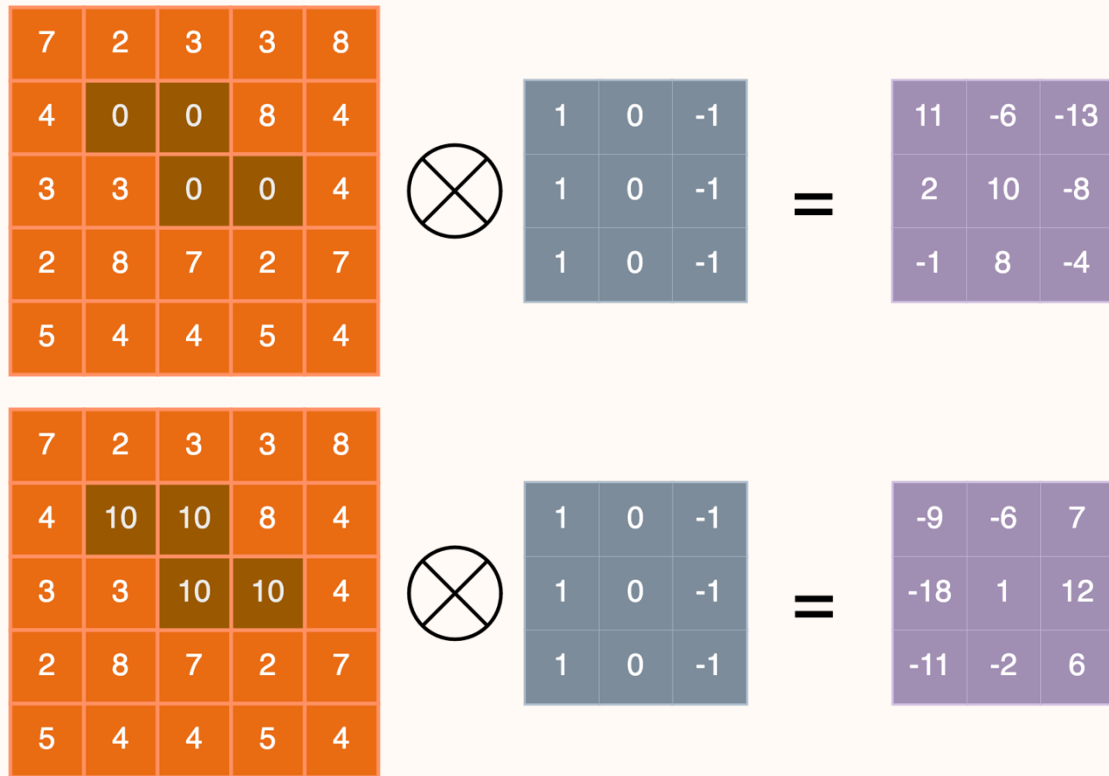


1	0	-1
1	0	-1
1	0	-1

=

6	-9	-8
-3	-2	-3
-3	0	-2

Convolutional layer



Partial Convolutions

- Only the valid pixels contribute to the convolution

$$x' = \begin{cases} \mathbf{W}^T (\mathbf{X} \odot \mathbf{M}) \frac{\text{sum}(\mathbf{1})}{\text{sum}(\mathbf{M})} + b, & \text{if } \text{sum}(\mathbf{M}) > 0 \\ 0, & \text{otherwise} \end{cases}$$

Partial Convolutions

7	2	3	3	8
4	0	0	8	4
3	3	0	0	4
2	8	7	2	7
5	4	4	5	4



1	0	-1
1	0	-1
1	0	-1

=

16.5	-10.8	-19.5
3	18	-12
-1.1	10.3	-5.1

Partial Convolutions

- After each operation, the mask is updated to indicate which regions have been filled

$$m' = \begin{cases} 1, & \text{if } \text{sum}(\mathbf{M}) > 0 \\ 0, & \text{otherwise} \end{cases}$$

Partial Convolutions

1	1	1	1	1
1	0	0	1	1
1	1	0	0	1
1	1	1	1	1
1	1	1	1	1



1	1	1
1	1	1
1	1	1

Metrics

- MSE loss: measures the pixel-per-pixel similarity

$$MSE = \frac{1}{N} \sum_{i,j} (A_{i,j} - B_{i,j})^2$$

Metrics

- Dice Coefficient: measures the overlap between two sets

$$Dice = \frac{2 * |A \cap B|}{|A| + |B|}$$

Metrics

- Dice Coefficient: measures the overlap between two sets

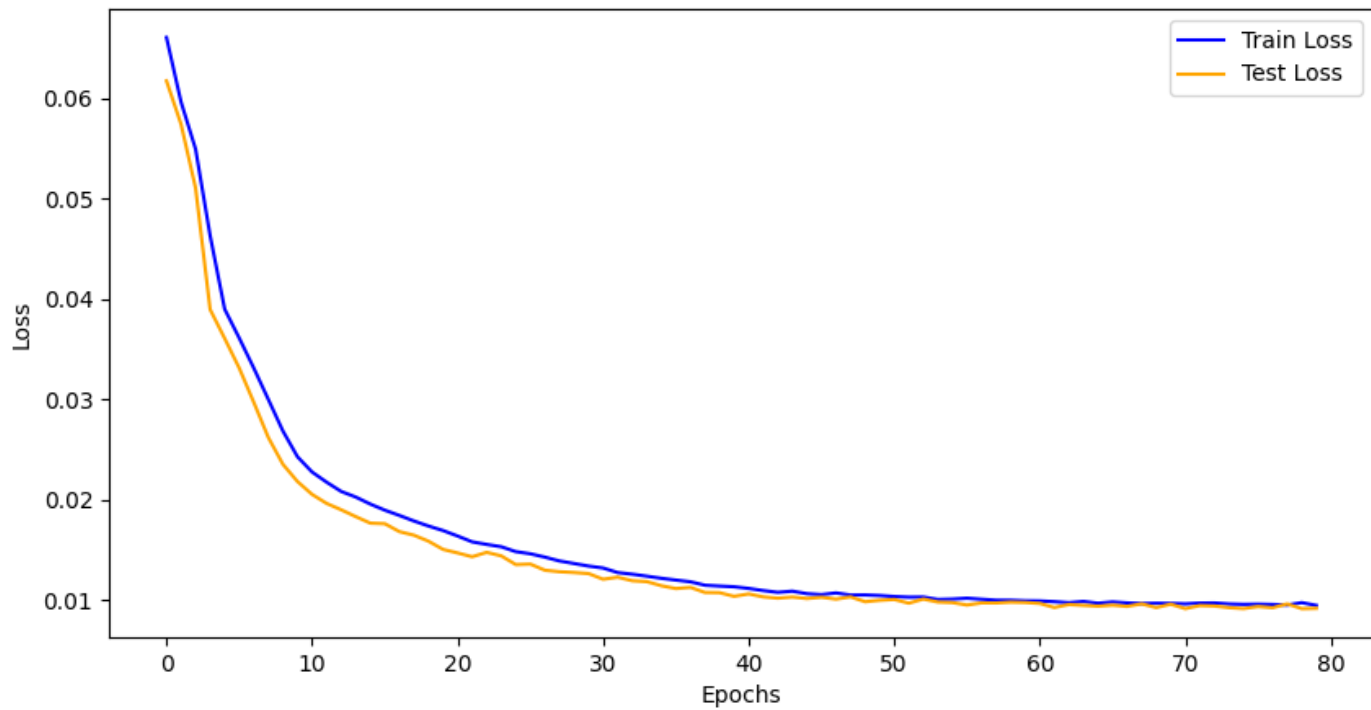
$$Dice = \frac{2 * \sum_{i,j} A_{i,j} B_{i,j}}{\sum_{i,j} A_{i,j} + \sum_{i,j} B_{i,j}}$$



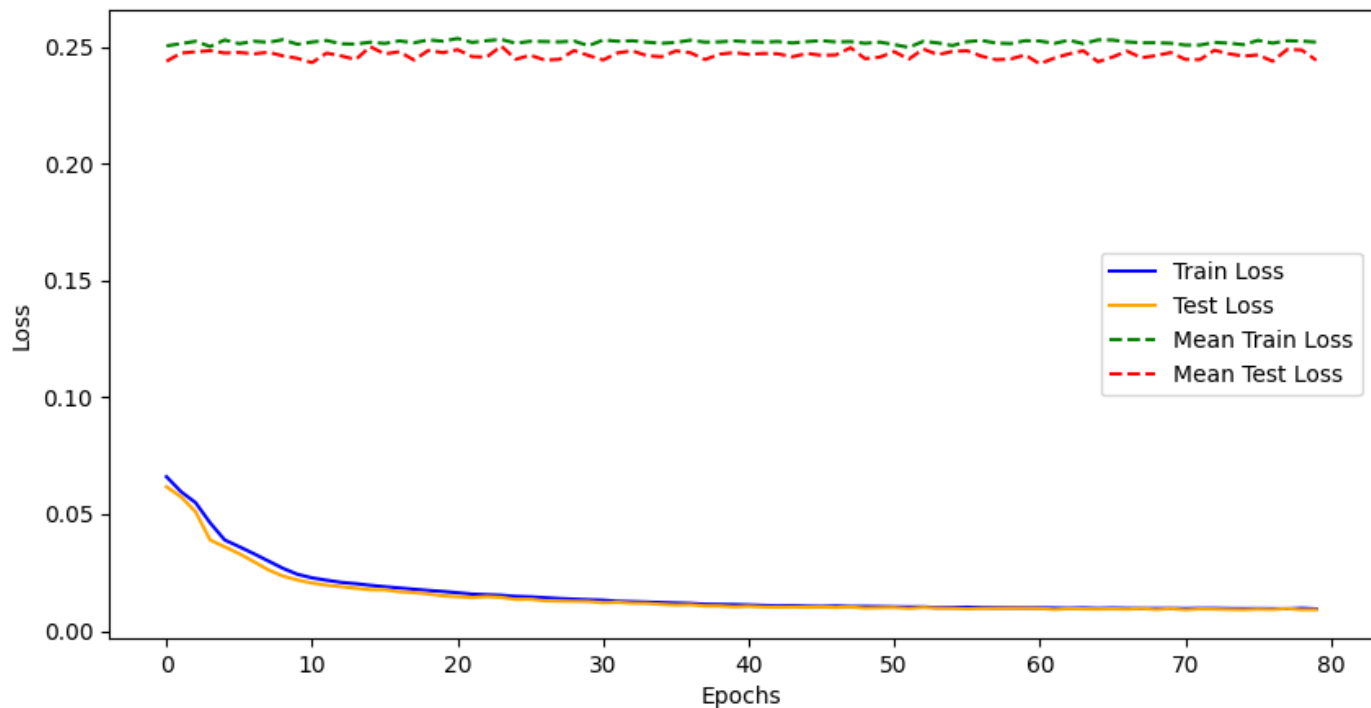
03

Results

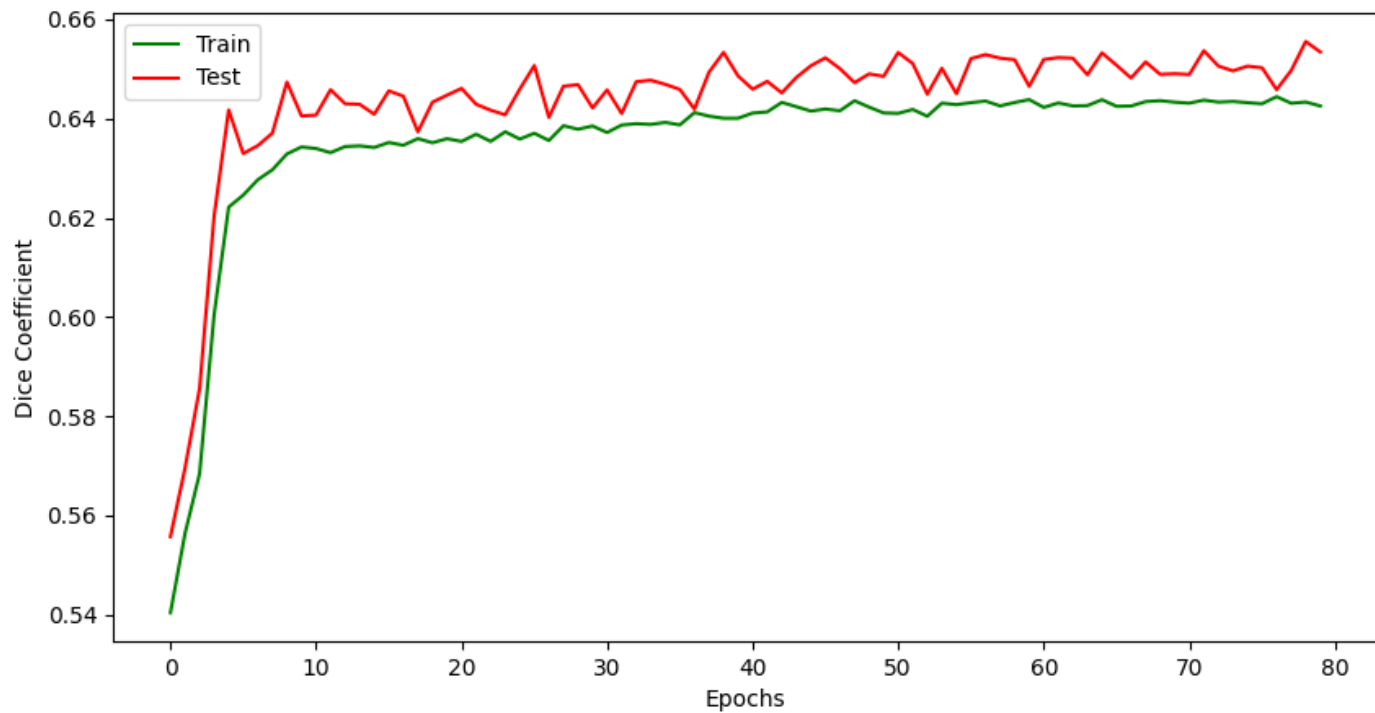
Loss



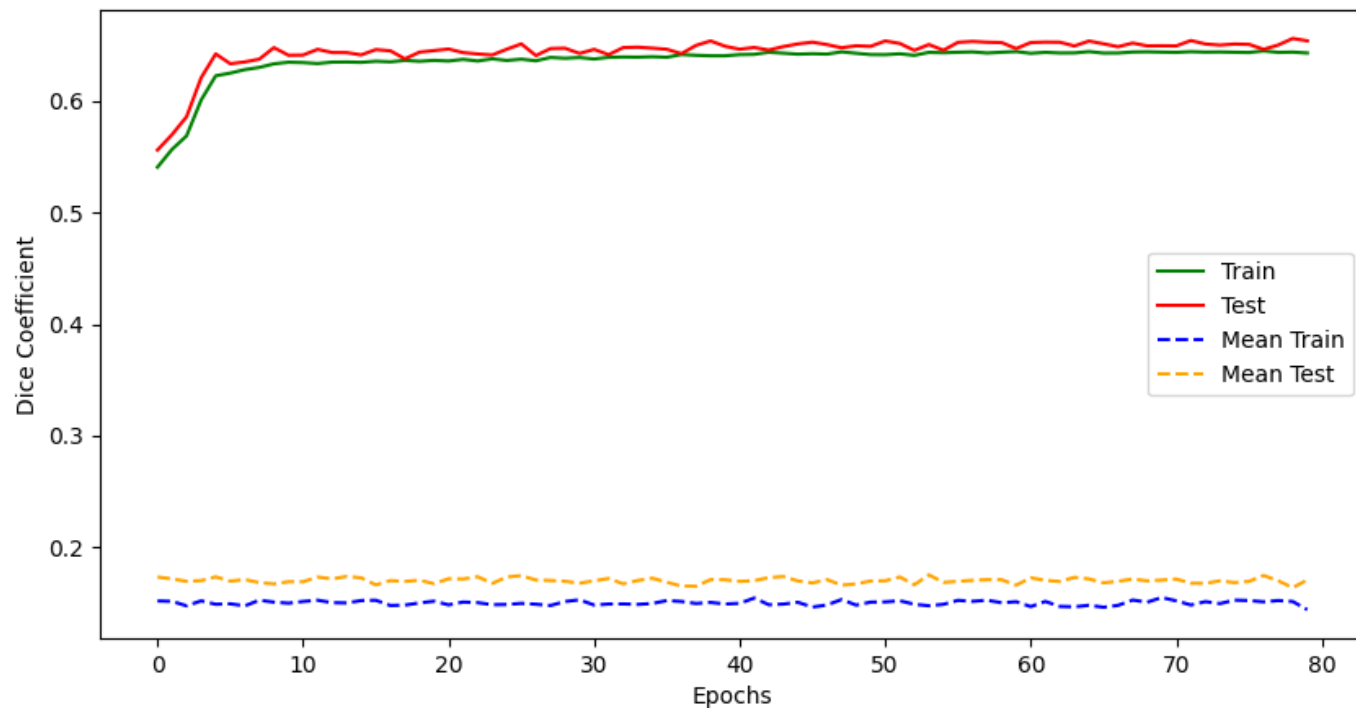
Comparison with the mean



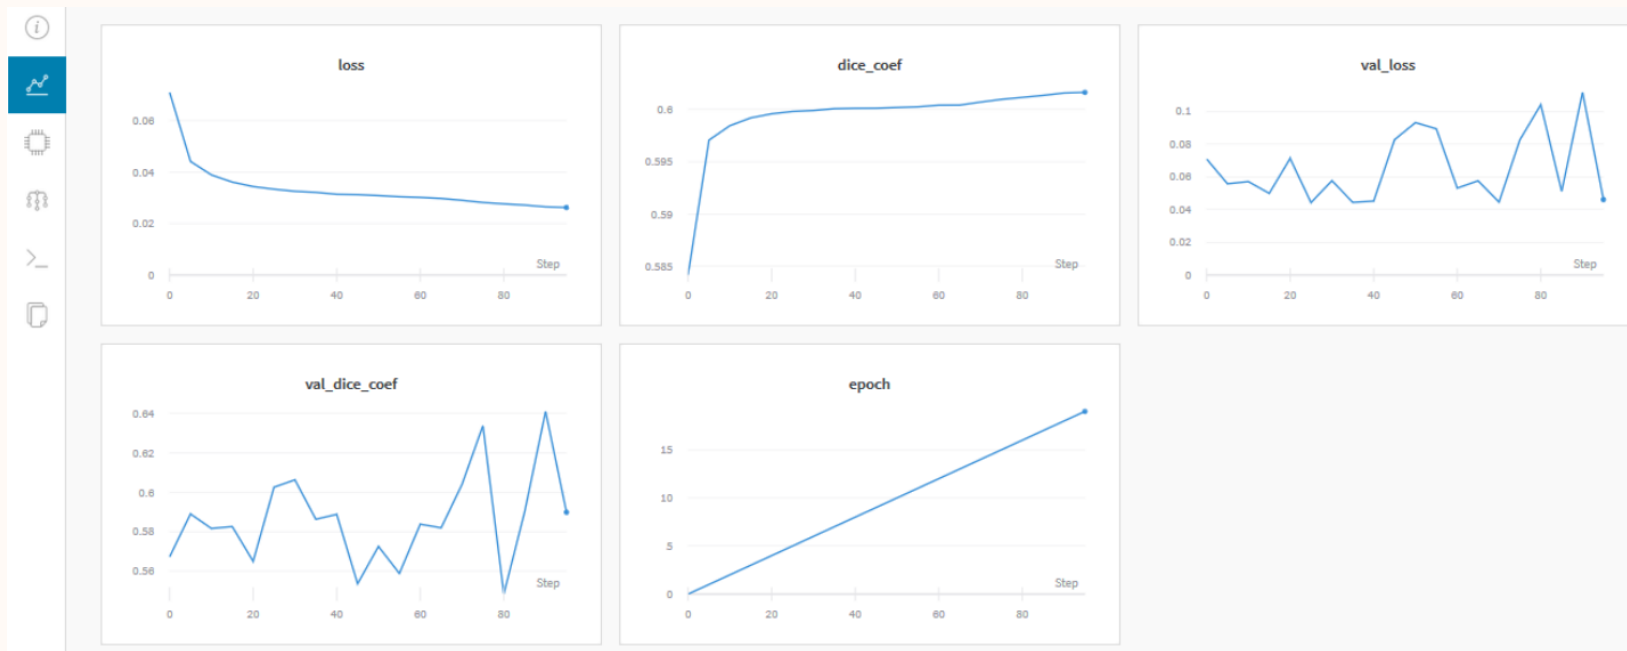
Dice



Comparison with the mean



Comparison



Hyperparameters

- Adam Optimizer
- MSE loss
- Initial learning rate 0.0002
- Learning rate multiplied by 0.1 every 10 epochs
- 50000 training images, 10000 test images
- Images normalized in the range $[-1, 1]$

Reconstructed Images

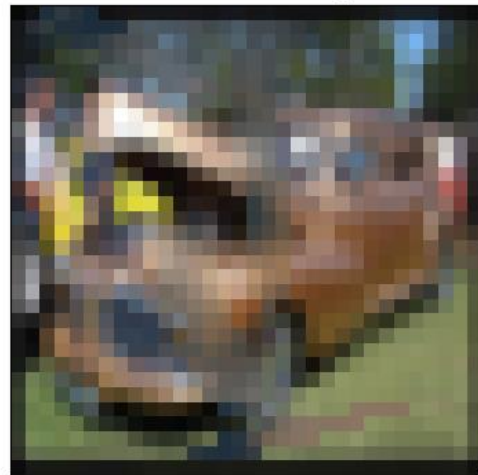
Original Image



Masked Image

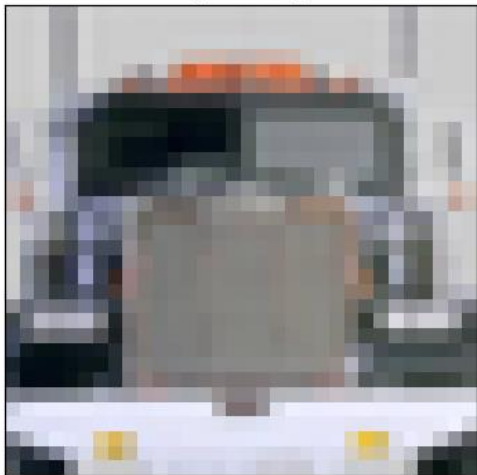


Reconstructed Image



Reconstructed Images

Original Image



Masked Image



Reconstructed Image



Reconstructed Images

Original Image



Masked Image

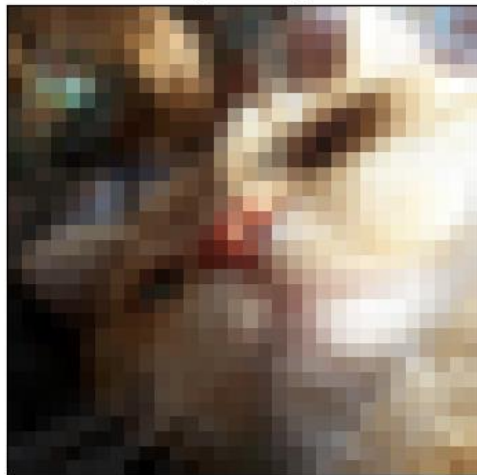


Reconstructed Image



Reconstructed Images

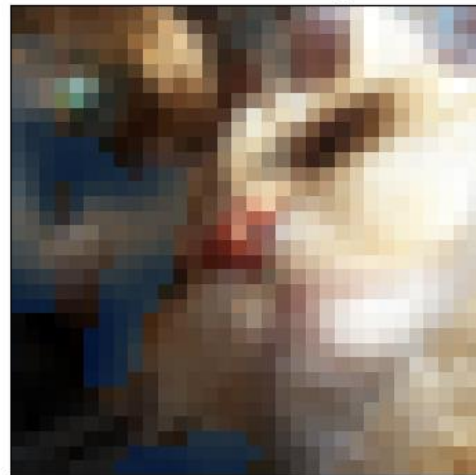
Original Image



Masked Image



Reconstructed Image





04

Improvements

More complex loss

From “Image Inpainting for Irregular Holes Using Partial Convolutions”, Liu et al.

$$\mathcal{L}_{total} = \mathcal{L}_{valid} + 6\mathcal{L}_{hole} + 0.05\mathcal{L}_{perceptual} + 120(\mathcal{L}_{style_{out}} + \mathcal{L}_{style_{comp}}) + 0.1\mathcal{L}_{tv}$$

Coarse-to-fine network

From “Generative Image Inpainting with Contextual Attention”, Yu et al.

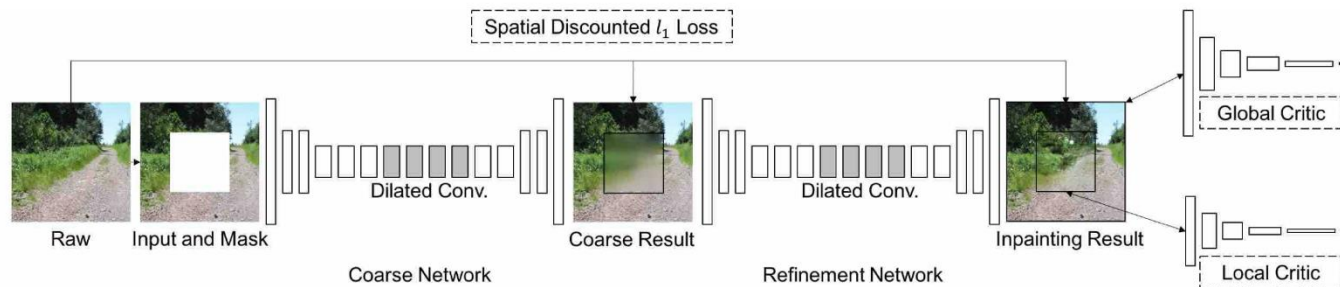


Figure 2: Overview of our improved generative inpainting framework. The coarse network is trained with reconstruction loss explicitly, while the refinement network is trained with reconstruction loss, global and local WGAN-GP adversarial loss.

Contextual Attention

From “Generative Image Inpainting with Contextual Attention”, Yu et al.

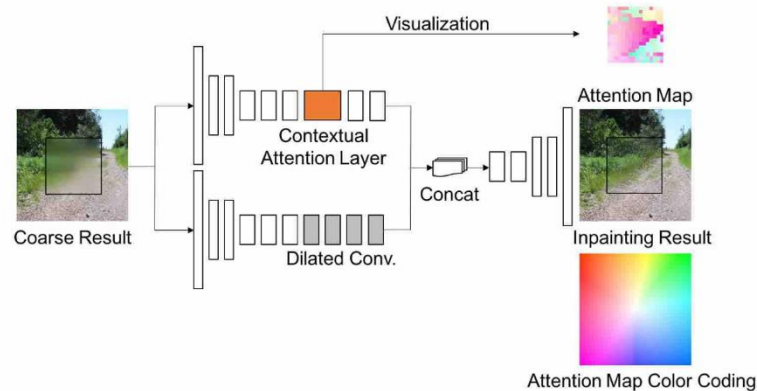


Figure 4: Based on coarse result from the first encoder-decoder network, two parallel encoders are introduced and then merged to single decoder to get inpainting result. For visualization of attention map, color indicates relative location of the most interested background patch for each pixel in foreground. For examples, white (center of color coding map) means the pixel attends on itself, pink on bottom-left, green means on top-right.



Thanks!

Sources:

- [Introduction to image inpainting with deep learning](#)
- [Generative Image Inpainting with Contextual Attention](#)
- [Image Inpainting for Irregular Holes Using Partial Convolutions](#)
- [deepimageinpainting](#) by ayulockin
- [PConv-Keras](#) by MathiasGruber



CREDITS: This presentation template was created by **Slidesgo**, and includes icons by **Flaticon** and infographics & images by **Freepik**

Thanks!